

Reetesh MUKUL  
Senior Principal Applied Scientist

in [linkedin.com/in/reetesh-mukul-5392532](https://www.linkedin.com/in/reetesh-mukul-5392532) [github.com/reeteshmukul](https://github.com/reeteshmukul)

IN : +91 89511 63698 @reetesh.mukul@gmail.com

Bangalore



I am currently working as a Senior Applied Principal Scientist in Computer Vision and Large Multi-Modality Models (LMM) domain. I specialize in developing large multi-modality models at various levels, including tasks specific to video and documents, as well as architectural modifications to incorporate multi-domain prompts and temporal aspects.

COMPÉTENCE

|                       |                                                                                                                                                                                                                                                                      |
|-----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Domain Expertise      | Deep Learning, Statistics, Computer Vision, Vision Transformers, Diffusion Models, Generative AI, Bayesian Programming, Advanced Mathematics, Classical Machine Learning, Parallel Programming, Reinforcement Learning, Functional Programming, Distributed Systems, |
| Programming Languages | C++20, Python, Haskell, Rust                                                                                                                                                                                                                                         |
| Frameworks            | BOOST, Torch, TIMM, OpenCV, PyRO, sklearn                                                                                                                                                                                                                            |
| Development Tools     | Visual Studio Code, Visual Studio, XCode, Git                                                                                                                                                                                                                        |
| Tech Pieces           | Transformers, CNN, Diffusion Models, Bayesian Programming, Pointer Networks, Linear Algebra, Batch Shop, Metaprogramming (C++), Statistical Programming                                                                                                              |

PROFESSIONAL EXPERIENCE

|          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|----------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Today    | Senior Principal Applied Scientist, ORACLE INDIA, Bangalore                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Dec 2023 | I am working on Large Multi-Modality Models with a focus on Document AI, Video Temporal Aspects, and Grounding. My work involves Model Training and Architectural developments.<br><div>LMM Transformers</div>                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Dec 2023 | Principle Machine Learning Engineer, SINGAPORE RESEARCH CENTER, HUAWEI, Singapore                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| Dec 2022 | My role is to conduct and coordinate Research Initiatives associated with Ad Creatives – understanding the content of the Advertisements and the cognitive interactions between various Ad Elements and users. The area thus involves Computational Advertisement, Creatives Psychology, Computer Vision, Cognitive analysis and Bayesian Programming. Most of my work requires developing Deep Learning Vision Models, Text Attention Models involving Advertisements. I work extensively on Neural Architectures, Image Attention, Transformers, Saliency and Memorability.<br><div>Torch Bayesian Programming Vision Transformers CLIP Denoising Diffusion</div> |

|          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|----------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Nov 2022 | <b>Principle Machine Learning Engineer, INDIA RESEARCH CENTER, HUAWEI, India</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| May 2021 | <ul style="list-style-type: none"> <li>➤ Ad Creative – Deep Learning-Driven Advertisement Creation. I work on Automatic Advertisement Creation; and Creative Ranking using various compositional aspects involving Advertisements. All factors, like – Industry, Product, Market, Users, Arts; – are involved. The system involved learning the causal impact of each component on the overall performance. My work further involves several CNN/Transformer based tools for Visual Analysis, Cognitive Analysis, Bayesian Programming for Reasoning, a few NLP and Sequential tools for learning the text impact, and Graphical Neural Network for understanding the connection among components.</li> <li>➤ Saliency Prediction – Network Development to predict Saliency</li> <li>➤ Blur Prediction – Network Development to predict Blur</li> <li>➤ Memorability – Network Development to predict Memorability</li> </ul> <p>Torch Advertisement VisionTransformers Memorability Cognition Denoising Diffusion</p> |
| Apr 2021 | <b>Machine Learning Developer, ADOBE, India</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| Apr 2020 | <ul style="list-style-type: none"> <li>➤ Sensei : ML Model development and deployment</li> <li>➤ Visual Similarity Model Development</li> <li>➤ Panoptic Segmentation : Module development on Post Processing side.</li> <li>➤ Metric development for ML models.</li> <li>➤ Tensor Library development for a subset of PyTorch</li> </ul> <p>Torch ONNX CoreML Keras Tensorflow CNN LSTM Clustering</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| Mar 2020 | <b>Senior Software Developer   Machine Learning Engineer, ADOBE, India</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| Sep 2016 | <ul style="list-style-type: none"> <li>➤ Lightroom Performance : Performance enhancement for Lightroom Classic Import, Grid, and Library</li> <li>➤ Markov Decision Process-based Dynamic Algorithm for Import : Parallel Distributed Batch Allocation based on MDP.</li> <li>➤ Performance Database for Lightroom : Light-weight embedded Performance Database that Profiles Lua code automatically.</li> <li>➤ Common Table Expressions for folders, keywords. Facilitates very fast lightweight low latency queries.</li> <li>➤ Feature Prediction : Natural Language Model for Photographic Features. Predicts future features based on previous features.</li> <li>➤ Analytics for Lightroom.</li> <li>➤ ML Models for Churn analysis and Garbage Collection</li> </ul> <p>XCode Visual Studio Lua SQLite R Python Keras Torch Hadoop Hive LSTM<br/>Reinforcement Learning Distributed Programming Scheduling Markov Decision Process</p>                                                                         |
| Aug 2016 | <b>Senior Software Developer   Machine Learning Engineer, FLIPKART, India</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| Aug 2015 | <ul style="list-style-type: none"> <li>➤ ML Model for Ad Click Prediction : CTR, CVR prediction and Ad Ranking</li> <li>➤ Ad SDK for Mobile Platform</li> </ul> <p>Ads Native Ads Pandas R Ranking Sampling Android Thompson Sampling<br/>Distributed Programming</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| Jul 2015 | <b>Senior Software Developer   Machine Learning Engineer, QUALCOMM, India</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| Dec 2011 | <ul style="list-style-type: none"> <li>➤ Optical Character Recognition for Indic (Devnagiri) Letters</li> <li>➤ Online detection of interesting scenes in Camera</li> </ul> <p>Computer Vision C++ OpenCV Android Python Assembly Parallel Programming</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |

|           |                                                                                             |
|-----------|---------------------------------------------------------------------------------------------|
| Nov 2011  | Software Developer, QUALCOMM, India                                                         |
| June 2008 | > Voice State Machine and Driver Development<br>C++ Embedded Systems Parallel Programming   |
| May 2008  | Software Developer, TEXAS INSTRUMENTS, India                                                |
| June 2006 | > Global Positioning Systems software development<br>C++ Embedded Systems Assembly Language |

## PROJECTS

---

Automatic Advertisement Creation keeping Ad Composition and Consumer Psychology in perspective

Creative Advertisement is the process and system for Creating Advertisements in an automated way. We study the dynamics of Market, Industry, Product, User, and Arts to create Advertisements. The created Advertisements are supposed to influence the Emotions and Psychology of the User so that there is a positive impact on the Product and the Brand.

Torch Congitive Sciences Bayesian Programming Computer Vision

### LRSENSEI (ADOBE)

2020 - 2021

Computer Vision Project Ranging From Best Photos, Auto Stacking to Panoptic Segmentation

Lightroom Sensei is a multiple-goal Computer Vision Project. I was working on Best Photos, Visual Similarity, and Panoptic Segmentation. My goal is to enhance and develop models, look for its deployment, try to find out suitable metrics, establish requirements by Users, and look for guarantees that different models provide. This work also involves choosing the deployment choices for the model -like can we put on client-side or on the server, how effectively jobs can be distributed on CPUs and GPUs.

Torch ONNX CoreML Feature Pyramids Contrastive Learning

### LRPERFORMANCE (ADOBE)

2016 - 2020

Performance Enhancement and User Analysis for Lightroom Classic

Lightroom Performance Enhancement had four major directions -(a) Efficient Algorithms, (b) Deep Parallelism, (c) Better Resource management, and (d) Machine Learning to achieve goals associated with (a)-(c). In an adversarial environment where User requirements are unknown, where Computation is costly (as Lightroom is a Photography Application) and Operations are done on multiple assets at the same time(here assets are Images), the challenges become multi-faceted. We successfully used Common Table Expressions of Database and System Parallelism to gain Performance. Thereafter we deployed Classical Reinforcement Learning based solutions, Markov Decision Process and Bandits to estimate resources. This is a very new paradigm, which has been successfully productized. Some of the developments, for example for Grid, involved estimating Geometry shifts using Weibull Distribution. Other than Performance, we used User Log data to predict future features, User churn rate, Probability of user to convert. We even tried to do Garbage Collection in Lua using Reinforcement Learning.

C++ Lua SQLite Markov Decision Process Reinforcement Learning Keras LSTM Hive Hadoop

### DEVNAGIRI OCR (QUALCOMM)

2011 - 2015

Optical Character Recognition for Devnagiri

. We developed Optical Character Recognizer for Indic (Devnagiri) Characters. We developed both Char Decoders and Word Decoders. This solution was developed for Mobile Devices hence challenges were on performance side as well. One of my goals was to develop cache aware vector routines. I developed Algorithm for Chandrabindu development and also wrote modules for Word Decoders.

C++ Computer Vision Machine Learning Android

## PAPERS AND PATENTS

1. Photo-editing application, Patent number : 10884769
2. Markov Decision Process for Efficient Data Transfer

3. Context-based Recommendation System for Feature Search
4. Environment Aware Application-based Resource Management Using Reinforcement Learning
5. Automatic Teeth Whitening Using Teeth Region Detection And Individual Tooth Location , United States Patent Application 20200342586
6. Photo-Editing Application Recommendations
7. Parameter Estimation for Accelerometers, Processes, Circuits, Devices and Systems. <https://patents.google.com/patent/US20090259424A1/en>
8. Broadband Hf/Vhf/Uhf Communication on Power Lines. <https://patents.google.com/patent/WO2007000777A1/de>.
9. R. Mukul et al., "An adaptive bandwidth request mechanism for QoS enhancement in WiMax real time communication," 2006 IFIP International Conference on Wireless and Optical Communications Networks, Bangalore, 2006, pp. 5 pp.-5, doi : 10.1109/WOCN.2006.1666583.

## EDUCATION

---

2006 **M.Tech** [IIITB Bangalore]

2004 **B.Tech** [BIT Sindri]

## LANGUAGES

---

English ● ● ● ● ●  
Hindi ● ● ● ● ●

## FORCES

---

- > Learner
- > Passionate
- > Motivated
- > Autonomus